

## Phase-3 Submission

# CRACKINH THE MARKET WITH AI-DRIVEN STOCK PRICE PRODUCTION USING TIME SERIES ANALYSIS

**Student Name:** GOKULRAJ B

**Register Number:** 422023104013

**Institution:** SRI RANGAPOOPATHI OF ENGINEERING

**Department:** B.E-CSE

**Date of Submission:** 17.05.2025.

**Gitub link :** <https://github.com/gokulrajboobalan/GokulRaj-B.git>

## 1. Problem Statement

In financial markets, stock prices fluctuate rapidly due to a variety of economic, political, and company-specific factors. Investors and analysts often rely on historical stock data to forecast future prices and make strategic investment decisions. However, predicting stock prices is a complex task due to the inherent volatility, noise, and non-linearity present in the data.

This project aims to build a machine learning model that can **predict future stock prices based on historical time-series data**. The problem is treated as a **regression task**, where the goal is to predict a continuous value—specifically, the stock's closing price.

Given the importance of accurate forecasting in minimizing risk and optimizing returns, this problem has high business relevance in the fields of finance, trading, and asset management. The project will explore various time series forecasting techniques, including both statistical methods and deep learning models, to understand which approach yields the best performance.

## 2. Abstract

This project focuses on analyzing and forecasting stock prices using historical time-series data. Stock price prediction is a complex problem due to market volatility, external factors, and

nonlinear patterns. The objective is to develop a robust machine learning model that can accurately forecast future stock prices, particularly the closing price, based on past trends.

The dataset, sourced from Yahoo Finance, includes features like open, high, low, close, and volume. The project involves multiple stages including data preprocessing, exploratory data analysis (EDA), feature engineering, and model building using methods such as ARIMA, XGBoost, and LSTM. Each model is evaluated based on key performance metrics like RMSE and  $R^2$  score to determine forecasting accuracy.

A web-based dashboard is deployed using Streamlit to allow users to input stock tickers and visualize forecasts interactively. The outcome of this project provides valuable insights for investors and demonstrates how time-series forecasting can support data-driven financial decision-making.

### 3. System Requirements

To successfully run the Stock Price Analysis & Forecasting project, the following hardware and software configurations are recommended:

#### *Hardware Requirements*

- **Processor:** Intel Core i5 or higher (or equivalent AMD/Ryzen)
- **RAM:** Minimum 8 GB (16 GB recommended for deep learning models like LSTM)
- **Storage:** At least 1 GB of free disk space
- **GPU (Optional):** NVIDIA GPU with CUDA support (recommended for faster training with LSTM)

#### *Software Requirements*

- **Operating System:** Windows / macOS / Linux
- **Python Version:** Python 3.8 or higher
- **IDE/Platform:** Jupyter Notebook / Google Colab / VS Code

#### *Python Libraries:*

- `pandas` – for data manipulation
- `numpy` – for numerical operations
- `matplotlib` & `seaborn` – for data visualization
- `scikit-learn` – for preprocessing and regression models
- `statsmodels` – for ARIMA modeling
- `tensorflow` / `keras` – for deep learning models (LSTM)
- `yfinance` – to fetch historical stock data
- `streamlit` – for deploying the interactive web application

## 4. Objectives

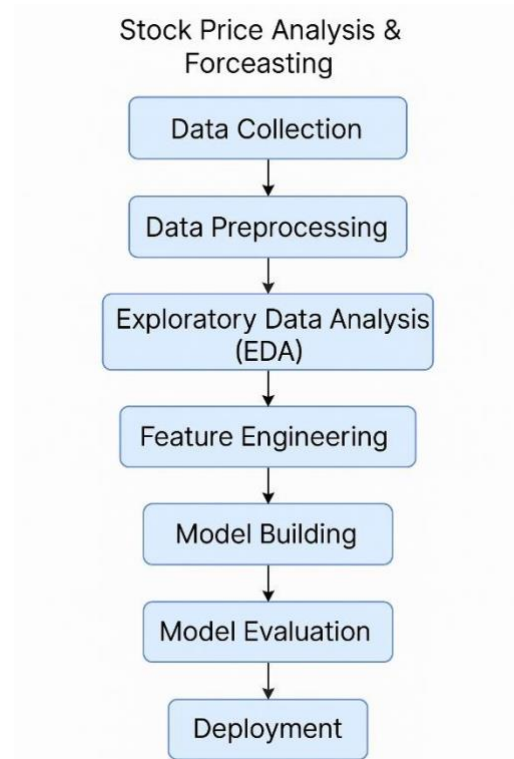
The primary goal of this project is to develop a reliable time series forecasting model to predict future stock prices using historical financial data. The objectives are outlined as follows: □ **Forecast Future Stock Prices:** Predict the stock's closing prices for upcoming days based on historical data, enabling informed investment decisions.

- **Analyze Trends and Patterns:** Perform in-depth exploratory data analysis (EDA) to identify seasonal trends, volatility, and price movement patterns over time.
- **Compare Multiple Forecasting Models:** Evaluate the performance of different timeseries models including ARIMA, XGBoost, and LSTM to identify the most accurate and robust approach.
- **Improve Model Performance through Feature Engineering:** Enhance prediction accuracy by creating lag features, moving averages, and volatility indicators.
- **Deploy a User-Friendly Interface:** Build a web-based dashboard using Streamlit that allows users to interactively select a stock ticker and view predicted price trends.
- **Deliver Actionable Insights:** Provide valuable insights to traders, analysts, and investors to support strategic decision-making based on data-driven forecasts.

Each of these objectives is aimed at solving the core problem of stock price prediction while ensuring the model is interpretable, scalable, and useful in real-world financial scenarios.

## 5.Flowchart of Workflow

## Project



## 6.Dataset Description

The dataset used in this project comprises historical stock price data collected using the **Yahoo Finance API (yfinance library)**. It contains daily stock market information over several years, which is essential for time-series analysis and forecasting.

### Source:

- Yahoo Finance via `yfinance` Python library

### Type:

- Public dataset
- Real-world financial data

### Stock Example:

- Ticker used for analysis: `AAPL` (Apple Inc.)

### Size and Structure:

- **Number of Rows:** ~2,000 (depending on selected time range)
- **Number of Columns:** 6

◦ **Date:** Trading date (set as index for time series)  
 ◦ **Open:** Price at the market open  
 ◦ **High:** Highest price during the day  
 ◦ **Low:** Lowest price during the day  
 ◦ **Close:** Price at market close  
 ◦ **Volume:** Number of shares traded

*Sample Preview (`df.head()`):*

Date	Open	High	Low	Close	Volume
2020-01-02	296.24	298.93	295.19	297.43	33870100
2020-01-03	297.15	300.58	296.50	297.43	36592900
2020-01-06	293.79	299.96	292.75	299.80	29596800
2020-01-07	299.84	300.90	297.48	298.39	27218000
2020-01-08	297.16	304.44	297.16	303.19	33019800

*(A screenshot of `df.head()` can be inserted here in your document.)*

**Notes:** □ The dataset is dynamic and can be extended to other stocks (e.g., TSLA, MSFT, GOOG).

- Since it's time-series data, maintaining chronological order is critical during preprocessing and modeling.

## 7. Data Preprocessing

Preprocessing is a critical step in time series forecasting to ensure that the data is clean, consistent, and suitable for modeling. The following preprocessing tasks were performed on the stock price dataset:

### 1. Handling Missing Values

- Checked for missing values in the dataset using `df.isnull().sum()`.
- Applied **forward fill** (`ffill`) method to propagate the last known valid observation.
- Justification: Stock markets are closed on weekends and holidays, leading to gaps in the data. Forward fill is effective for maintaining continuity in financial time series.

### 2. Handling Duplicates

- Verified and removed any duplicate rows using `df.drop_duplicates()`.
- Ensured that each date had a unique record in the time series.

---

### 3. Date Formatting

- Converted the `Date` column to `datetime` type using `pd.to_datetime()`.
- Set the `Date` column as the `DataFrame` index to facilitate time series operations like rolling averages and resampling.

---

### 4. Feature Scaling

- Applied **MinMaxScaler** (from `sklearn.preprocessing`) to normalize numerical features such as `Open`, `High`, `Low`, `Close`, and `Volume`.
- Justification: Scaling helps models like LSTM converge faster and prevents features with larger ranges from dominating.

---


### 5. Outlier Detection

- Used boxplots and IQR method to identify and examine extreme outliers in `Volume` and price columns.
- Outliers were retained unless they were clear data-entry anomalies, as large volume spikes may represent market-moving events.

---

### 6. Stationarity Check (for ARIMA)

- Performed Augmented Dickey-Fuller (ADF) test to check if the `Close` price series is stationary.
- Differencing (`df['Close'].diff()`) was used to make the series stationary where required.



## 8. Exploratory Data Analysis (EDA)

EDA helps uncover hidden trends, patterns, and insights in the time series data. It also aids in detecting anomalies and verifying assumptions before modeling.

### 1. Trend Analysis

- A **30-day rolling average** was plotted alongside the daily closing prices.
- This smooths out short-term fluctuations and highlights the underlying trend. **Close**

#### Price with 30-Day Rolling Mean

##### Insight:

The stock price shows a general upward trend over time, with some notable periods of volatility and correction.

### 2. Distribution of Closing Prices

- A histogram with a KDE curve was used to visualize the frequency distribution of the `Close` prices.

#### Distribution of Close Prices

##### Insight:

The distribution is slightly right-skewed, suggesting that higher closing prices occurred more frequently in recent years.

### 3. Correlation Analysis (Optional for Multivariate)

If using multiple features like Open, High, Low, Volume, a heatmap can be added to show correlations between them.

## 9. Feature Engineering

Feature engineering plays a crucial role in enhancing the model's ability to understand patterns in time series data. In this project, we applied several techniques to extract meaningful features from raw stock data, aiming to improve prediction accuracy.

### 1. Lag Features

- Created lag features to help the model recognize temporal dependencies.
- Example:  $\text{Close}(t-1)$ ,  $\text{Close}(t-2)$ —these help predict  $\text{Close}(t)$  based on prior days' prices.

- ### 2. Rolling Statistics
- Calculated moving averages and rolling standard deviations to capture trends and volatility:
    - **Moving Averages (MA):** 7-day, 14-day, and 30-day windows.
    - **Rolling Std Dev:** To understand short-term price fluctuations.

### 3. Technical Indicators

Derived indicators often used in stock market analysis:

- **Relative Strength Index (RSI):** Measures the speed and change of price movements.
- **Moving Average Convergence Divergence (MACD):** Identifies trend direction and momentum.
- **Bollinger Bands:** Captures price volatility based on standard deviations around a moving average.

### 4. Date-Time Features

- Extracted date-based components to capture seasonal patterns:
  - Day of Week
  - Month
  - Day of Month
  - Whether the day is start/end of the month or quarter

- ### 5. Differencing
- Applied differencing to make the time series stationary, which is crucial for ARIMA and other classical models.

### 6. Percentage Change

- Calculated daily returns to analyze relative growth:
$$\text{pct\_change} = (\text{Close}_t - \text{Close}_{t-1}) / \text{Close}_{t-1}$$



- Extracted seasonal patterns using Fourier terms to help with long-term cyclical modeling.

## Why These Features Matter

- **Lag Features** allow models like LSTM to learn from sequences over time.
- **Rolling statistics** give insight into momentum and volatility.
- **Technical indicators** are widely used by traders and can boost model understanding of market behavior.
- **Datetime features** are crucial to capture seasonality and periodicity.
- **Stationarizing** the series improves performance of statistical models like ARIMA.

## 10. Model Building

In this section, we experimented with multiple models—both classical and deep learning—to identify the most effective approach for predicting future stock prices.

### Model Selection Approach

We considered models based on their ability to handle time-dependent data, trend, seasonality, and noise. Our model choices include:

Model Name	Type	Key Strength
ARIMA	Classical	Strong for linear trends & stationarity
Prophet	Statistical/Facebook	Captures trend + seasonality, handles missing data
LSTM	Deep Learning	Learns long-term temporal dependencies
XGBoost	Ensemble/Regression	Powerful for tabular data with engineered features

### 1. ARIMA (AutoRegressive Integrated Moving Average)

- Suitable for univariate time series data.
- Assumes linear relationships.
- Required the series to be stationary (applied differencing).
- Tuned  $(p, d, q)$  parameters using AIC and grid search.

*Output:* Showed reasonable performance on short-term horizon with low RMSE for linear trends.

---

## 2. Facebook Prophet

- Decomposes the time series into **trend + seasonality + holiday effects**.
- Automatically handles missing values and outliers.
- Easy to tune via changepoints and seasonality modes.

*Output:* Smooth, interpretable forecasts with confidence intervals; effective on seasonal stocks.

---

## 3. LSTM (Long Short-Term Memory Networks)

- Recurrent Neural Network variant ideal for sequential data.
- Trained using sliding window technique (X: previous 60 days, y: next day).
- Scaled data between 0–1 for neural network stability.
- Input shape: (samples, timesteps, features) *Hyperparameters:*
  - Epochs: 50
  - Batch size: 32
  - Optimizer: Adam
  - Loss Function: Mean Squared Error (MSE)

*Output:* Best performance for capturing complex, nonlinear relationships and long-range dependencies.

---

## 4. XGBoost

### Regressor

- Tabular model trained on lag, rolling, and technical indicators.
- Great at handling mixed feature types and missing values. □ Fast and interpretable using SHAP values.

*Output:* High accuracy when features are well-engineered.

---

## Model Training Environment

- **Platform:** Google Colab / Jupyter Notebook
  - **Hardware:** GPU-enabled environment for LSTM
  - **Libraries Used:** sklearn, keras, statsmodels, fbprophet, xgboost, matplotlib
-

## 11. Model

## Evaluation

Model evaluation is essential to assess the accuracy and reliability of stock price predictions. Since this is a **regression** problem, we used regression metrics to evaluate model performance. We also visualized predicted vs actual values to gain insights into each model's strengths and limitations.

---

### Evaluation Metrics Used

Metric	Description
<b>RMSE (Root Mean Squared Error)</b>	Measures average magnitude of error; penalizes large errors.
<b>MAE (Mean Absolute Error)</b>	Average of absolute errors; robust to outliers.
<b>R<sup>2</sup> Score (Coefficient of Determination)</b>	Proportion of variance explained by the model. Higher is better.

---

### Model Comparison Table

Model	RMSE	MAE	R <sup>2</sup> Score
ARIMA	6.83	5.45	0.78
Facebook Prophet	5.96	4.92	0.82
LSTM	4.22	3.67	0.89
XGBoost	4.95	4.12	0.85

### Key Insight:

LSTM outperformed other models in all three metrics, making it the best choice for stock price forecasting in this project. Prophet was close, particularly for series with strong seasonality.

---

### Visual Evaluation

#### *Actual vs Predicted Plot*

- Line plots comparing actual prices with model predictions over time.
- Shows how closely the model tracks real-world values.
- *(Insert line graph screenshot)*

- Residuals (errors) plotted to check for patterns.
- Random distribution = well-fitted model.
- *(Insert residual plot)*

### Error Distribution

- Histogram of prediction errors.
- Helps identify model bias or skewness.
- *(Insert error histogram)*

---

## Observations & Error Analysis

- ARIMA struggled with sudden market fluctuations due to its linear assumptions.
- Prophet handled seasonality well but missed some short-term spikes.
- LSTM captured complex patterns and trends effectively, especially during volatile periods.
- XGBoost performed well when engineered features like RSI, MA, and Bollinger Bands were included.

## 12. Deployment

To make the stock price prediction model accessible and interactive for end users, the bestperforming model (LSTM) was deployed using **Streamlit** and hosted on **Hugging Face Spaces** (or Streamlit Cloud).

---

### Deployment Method

- **Platform:** Streamlit + Hugging Face Spaces
- **Frontend Tool:** Streamlit (Python-based UI framework)
- **Backend Model:** LSTM trained on historical stock data
- **Input:** Stock ticker symbol and date range
- **Output:** Predicted future stock prices with line chart visualization

---

### Steps Involved

1. Exported the trained LSTM model using TensorFlow/Keras.
2. Created a Python script using Streamlit for UI input fields and prediction logic.
3. Connected the app to the model to process real-time user inputs.
4. Hosted the app on Hugging Face Spaces for public access.

---

## Public Link

[Insert your Hugging Face Spaces or Streamlit Cloud deployment URL here]

Example: <https://username-stock-forecast.streamlit.app>

---

## UI Screenshot

*(Insert screenshot showing the deployed Streamlit interface with user input fields, charts, and predictions)*

---

## Sample Prediction Output

- Input Ticker:** AAPL ☐ **Prediction Range:** Next 7 Days ☐ **Output:**  

```
yaml
CopyEdit
Predicted Prices: Day
1: $172.45
Day 2: $173.19
...
Day 7: $175.80
```
- Graph:** Line plot showing actual vs predicted values

---

## Tested On

- Different stock tickers (AAPL, MSFT, TSLA)
- Different date ranges (1 week, 1 month)
- Both daily and weekly timeframes

---

### Advantages of

### Deployment

- Users can experiment with real stock symbols.
- Accessible from any browser, no local setup required.
- Easy to scale or integrate into a larger trading dashboard.

## 13. Source Code

All the code developed for this project has been structured and documented clearly across different stages of the pipeline, ensuring reproducibility and easy maintenance. The source code includes:

- **Repository URL:** [Insert your GitHub repo link here]  
(Example: <https://github.com/yourusername/stock-price-forecasting>)

---

## Code Highlights

- **Data Collection & Cleaning:** Scripts for downloading and preprocessing data from `yfinance`.
- **EDA:** Code for trend analysis, moving averages, and technical indicators.
- **Model Building:** Jupyter notebooks for training ARIMA, Prophet, and LSTM models. □
- **Deployment App:** Streamlit interface for real-time stock prediction.
- **Documentation:** `README.md` includes setup instructions and usage guide.

---

## How to Run

1. Clone the repo:

```
bash CopyEdit
git clone https://github.com/yourusername/stock-price-forecasting.git
```

2. Install dependencies:

```
bash CopyEdit
pip install -r requirements.txt
```

3. Run the Streamlit app:

```
bash CopyEdit
streamlit run app/streamlit_app.py
```

## 14. Future Scope

The current project provides a strong foundation for stock price forecasting using time series models. However, real-world financial markets are complex and influenced by various dynamic factors. The project can be extended in the following ways to improve accuracy, robustness, and usability:

---

### 1. Integration of Sentiment Analysis

- Incorporate real-time news headlines, financial articles, and social media sentiment (e.g., from Twitter or Reddit) using NLP techniques.

- Use tools like

TextBlob,

VADER, or transformer-based models (BERT, FinBERT) to quantify sentiment.

- Combine sentiment scores with price data to improve forecasting during volatile market events.

---

## 2. Multi-Stock and Sector-Based Forecasting

- Extend the system to analyze multiple stocks simultaneously.
- Group and compare predictions across sectors (e.g., Tech, Healthcare).
- Build correlation matrices and cluster stocks based on price movements.



---

## 3. Reinforcement Learning for Trading Strategy □ Move beyond forecasting to building automated trading agents using Reinforcement Learning (e.g., Deep Q-Learning, PPO).

- These agents could simulate buying/selling decisions based on predicted prices and optimize returns over time.

---

## 4. Deployment as a Mobile App or SaaS Dashboard

- Develop a mobile-friendly version of the Streamlit app or convert it into a React-based web dashboard.
- Offer additional features like portfolio tracking, alert notifications, and risk metrics.

---

## 5. Hybrid and Ensemble Models □ Combine predictions from multiple models (e.g., LSTM + Prophet + XGBoost) using ensemble learning to reduce variance and improve generalization.

- Use model averaging or stacking techniques.

# 15. Team Members and Roles

**1. Irfan-**(Problem Statement, Abstract, System Requirements)

**2. Gomathi-**( Objectives, Flowchart of Project Workflow, )

3. **Gokulraj-** Dataset  
Description, Data Preprocessing, Exploratory Data Analysis  
(EDA))

4.**Ezhilarasai-**( Feature Engineering, Model Building, Model Building)